

Markov Chain Monte Carlo

Dustin Lang

2019-09-17

Markov Chain Monte Carlo for data analysis

- ▶ MCMC generates samples from a probability distribution function
- ▶ For data analysis: you have some *measurements* X , with uncertainties
- ▶ have a *model* for the thing you're measuring
- ▶ with unknown *parameters* θ
- ▶ the model is *generative* – predicts measurements given parameters
- ▶ MCMC can put constraints on the parameters given your measurements

Markov Chain Monte Carlo for data analysis

- ▶ can think of it as: *model* is a function mapping parameters to observables $t: m(\theta) \rightarrow t$
- ▶ and *measurement process* gives probability of observing X given t : $P(X|t)$
- ▶ (eg, maybe $P(X|t)$ is a Gaussian with zero mean and standard deviation σ)
- ▶ Then, can chain $P(X|t) = P(X|m(\theta))$
- ▶ Usually, we fold m and P together to write $P(X|\theta)$

Markov Chain Monte Carlo for data analysis

- ▶ $P(X|\theta)$ (the “likelihood”) is not what you want!
- ▶ You want $P(\theta|X)$ – the “posterior” – constraints on / measurements of the parameters
- ▶ Good thing you’re Bayesian!
- ▶ “posterior” $P(\theta|X) = \frac{P(X|\theta)P(\theta)}{P(X)}$
- ▶ $P(\theta)$ is the “prior”
- ▶ $P(X)$ is the “evidence”, and in this setting is constant (ignored)

Markov Chain Monte Carlo for data analysis

- ▶ We'll use MCMC to draw samples from the *posterior* probability distribution
- ▶ $P(\theta|X) = P(X|\theta)P(\theta)$
- ▶ The samples are a set of θ values drawn from your parameter constraints distribution

What's in a name

- ▶ *Monte Carlo* – a famous casino in Monaco – alluding to the use of *randomness* in the algorithm
- ▶ *Chain* – a list (here, a list of samples / steps $\theta_1, \theta_2, \theta_3$)
- ▶ *Markov Chain* – a list of “steps” where the decision to step from θ_i to θ_{i+1} depends only on the value of θ_i
- ▶ *Metropolis–Hastings* – first authors of 1953 and 1970 papers (resp) with the simplest, most commonly taught algorithm

Using MCMC

- ▶ “Just write down the probability distributions”
- ▶ hah, the “model to predictions” part, $m(\theta) \rightarrow t$, can be *very* difficult
- ▶ eg, in cosmology, the software *CAMB* does this, in 25,000 lines of FORTRAN
- ▶ writing down the *likelihood* might take some work too!

MCMC: the idea

- ▶ move a “particle” or “walker” around in the parameter space θ , based on the value of the probability function (for us: the posterior probability $P(\theta|X)$)
- ▶ choose new parameters according to a *proposal distribution* – often Gaussian: $\theta_{new} = \theta + \mathcal{N}(0, \Sigma)$
- ▶ compute probability at new θ_{new} and decide whether to move (“jump”) to new parameters
- ▶ *always* keep improvements, *sometimes* keep decreases
- ▶ in the long run, the set of walker positions $\{\theta_1, \theta_2, \theta_3, \dots\}$ are your samples


Finally, M–H MCMC code

```
fun MCMC(prob_function, sample_proposal, params, Nsteps):  
    chain = []  
    prob = prob_function(params)  
    for i in 1 to Nsteps:  
        # sample from proposal distribution  
        params_new = sample_proposal(params)  
        prob_new = prob_function(params_new)  
        # accept?  
        if ((prob_new > prob) OR  
            (prob_new / prob) > uniform_random()):  
            params = params_new  
            prob = prob_new  
        chain.append(params)  
    return chain
```


Demo

Credit: Dan Foreman–Mackey

<https://speakerdeck.com/dfm/data-analysis-with-mcmc>




Metropolis-Hastings





Metropolis-Hastings
in an ideal world


start here
perhaps



Metropolis-Hastings
in an ideal world




Metropolis-Hastings
in an ideal world



$$p(\text{accept}) = \min \left(1, \frac{p(\mathbf{x})}{p(\mathbf{x}')} \frac{q(\mathbf{x}; \mathbf{x}')}{q(\mathbf{x}'; \mathbf{x})} \right)$$


Metropolis-Hastings
in an ideal world



$$p(\text{accept}) = \min \left(1, \frac{p(\mathbf{x})}{p(\mathbf{x}')} \frac{q(\mathbf{x}; \mathbf{x}')}{q(\mathbf{x}'; \mathbf{x})} \right)$$

definitely.

Metropolis-Hastings
in an ideal world




$$\mathbf{x}' \sim q(\mathbf{x}'; \mathbf{x})$$

only **relative** probabilities


$$p(\text{accept}) = \min \left(1, \frac{p(\mathbf{x})}{p(\mathbf{x}')} \frac{q(\mathbf{x}; \mathbf{x}')}{q(\mathbf{x}'; \mathbf{x})} \right)$$

definitely.


Metropolis–Hastings
in an ideal world




Metropolis–Hastings
in an ideal world



Metropolis–Hastings
in an ideal world




Metropolis-Hastings
in an ideal world




not this time.


Metropolis-Hastings
in an ideal world




Metropolis–Hastings
in an ideal world




Metropolis-Hastings
in an ideal world



Metropolis-Hastings
in an ideal world




Metropolis-Hastings
in an ideal world




Metropolis-Hastings


in the real world




Metropolis–Hastings
in the real world




Metropolis-Hastings
in the real world



Metropolis-Hastings
in the real world




Metropolis–Hastings
in the real world




the
Small Acceptance Fraction
problem

Metropolis–Hastings
in the real world




Metropolis-Hastings

in the real world



Metropolis–Hastings
in the real world



the
Huge Acceptance Fraction
problem

Metropolis–Hastings
in the real world

MCMC: summary

Things MCMC can do for you

- ▶ produce constraints on model parameters, including covariances
- ▶ include (but ignore) *nuisance parameters* in your data analysis

Things MCMC doesn't do for you

- ▶ tell you if it has converged!
- ▶ tell you if your model is good
- ▶ (easily) allow comparisons between models
- ▶ handle multi-modal distributions (try *nested sampling*)
- ▶ Dan Foreman-Mackey has good notes on **advanced sampling methods**